

Project 5 – Recognition using Deep Networks

Abinav Anantharaman | NUID: 002774223

Satwik Shridhar Bhandiwad | NUID: 002920338

Introduction

This project focuses on building, training, analyzing, and modifying a deep neural network for the recognition of handwritten digits using the MNIST dataset. The goal is to demonstrate the potential of deep learning by creating a model that can accurately identify handwritten digits. The project is designed to be accessible without the need for a GPU and provides a good example of the capabilities of deep neural networks. Throughout the project, various techniques will be explored to improve the model's performance, providing insight into the process of building and improving deep learning models.

Task 1A: Get the MNIST digit data set

Ground Truth: 3




Ground Truth: 2



Ground Truth: 3




Ground Truth: 0



Ground Truth: 3




Ground Truth: 1



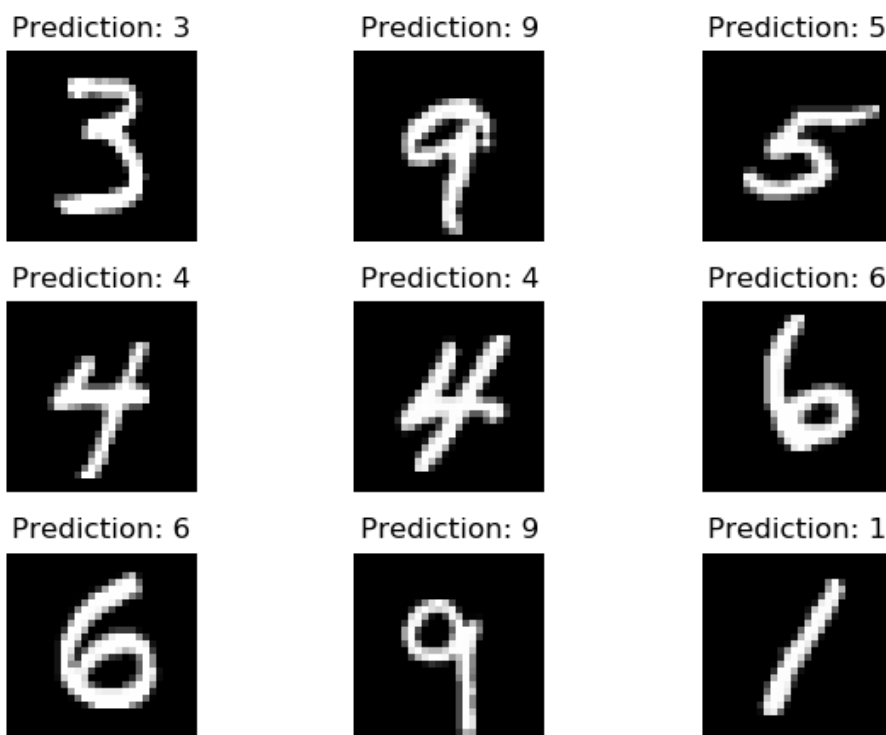
Task 1C: Build a network model

```
-----  
Layer (type)              Output Shape              Param #  
-----  
Conv2d-1                   [-1, 10, 24, 24]         260  
Conv2d-2                   [-1, 20, 8, 8]           5,020  
Dropout2d-3                [-1, 20, 8, 8]           0  
Linear-4                   [-1, 50]                 16,050  
Linear-5                   [-1, 10]                 510  
-----  
Total params: 21,840  
Trainable params: 21,840  
Non-trainable params: 0  
-----  
Input size (MB): 0.00  
Forward/backward pass size (MB): 0.06  
Params size (MB): 0.08  
Estimated Total Size (MB): 0.15  
-----  
Net(  
  (conv1): Conv2d(1, 10, kernel_size=(5, 5), stride=(1, 1))  
  (conv2): Conv2d(10, 20, kernel_size=(5, 5), stride=(1, 1))  
  (conv2_drop): Dropout2d(p=0.5, inplace=False)  
  (flatten): Flatten(start_dim=1, end_dim=-1)  
  (fc1): Linear(in_features=320, out_features=50, bias=True)  
  (fc2): Linear(in_features=50, out_features=10, bias=True)  
)
```

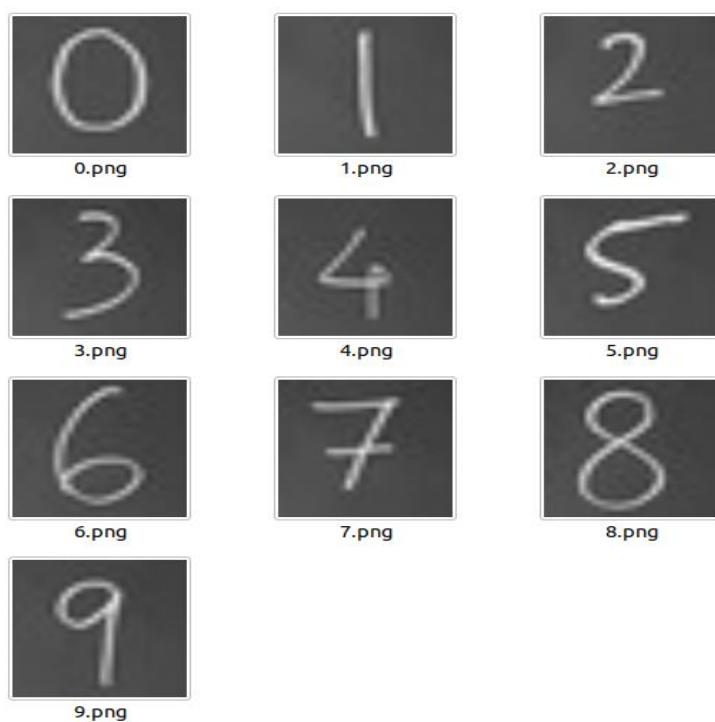
Task 1D: Train the model



Task 1F: Read the network and run it on the test set



Task 1G: Test the network on new inputs



Handwritten digits resized to 28 x 28 and converted to grayscale and inverted.

Prediction: 3



Prediction: 0



Prediction: 5



Prediction: 1



Prediction: 4



Prediction: 1



Prediction: 6



Prediction: 1



Prediction: 8



Prediction: 2




Got 80 percent accuracy on above handwritten image dataset.

Task 2A: Analyze the first layer


Filter 1




Filter 2




Filter 3




Filter 4



Filter 5



Filter 6




Filter 7



Filter 8



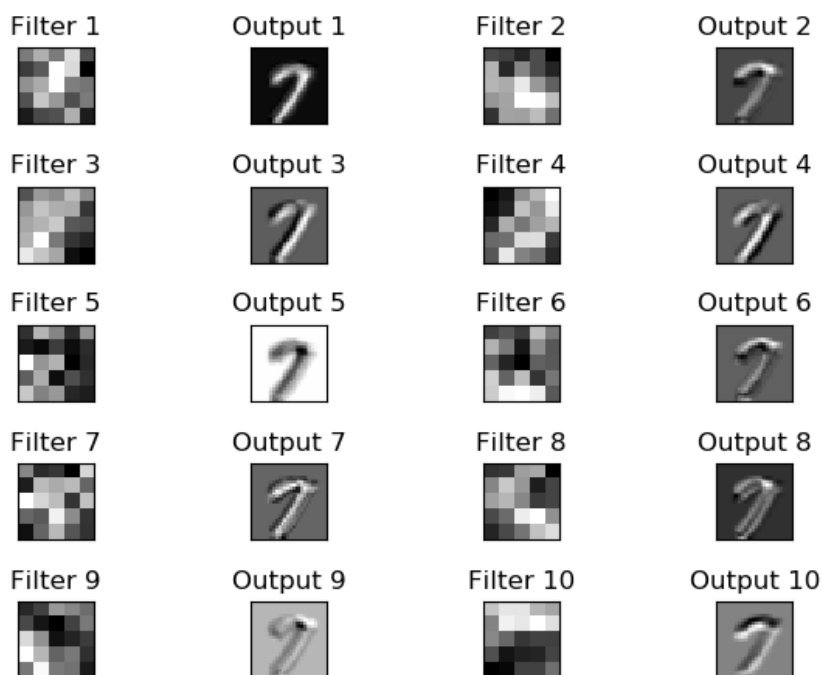
Filter 9



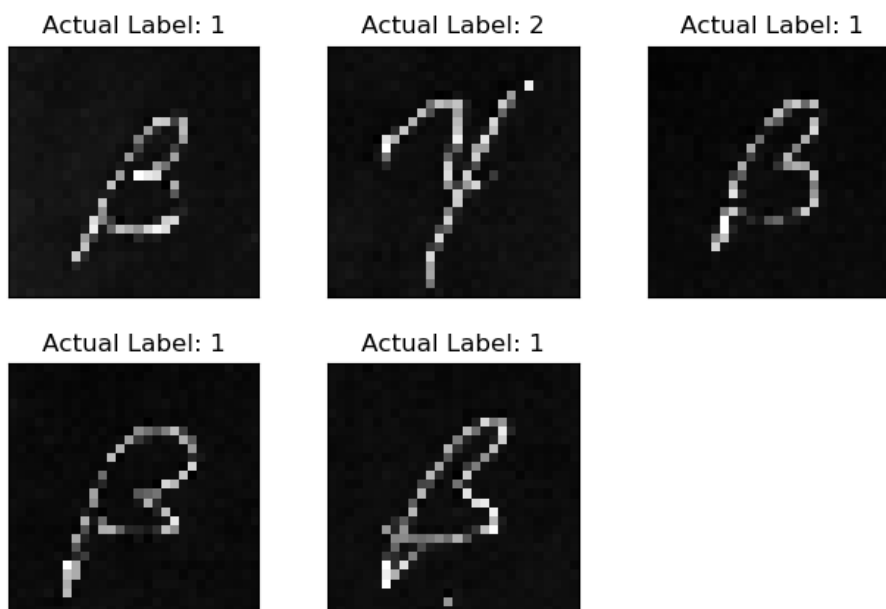
Filter 10




Task 2B: Show the effect of filters



Task 3: Transfer Learning on Greek Letters



Above are the sample images used for training



Above is the training Loss for Greek letters.
























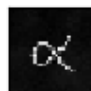



Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 10, 24, 24]	260
Conv2d-2	[-1, 20, 8, 8]	5,020
Dropout2d-3	[-1, 20, 8, 8]	0
Linear-4	[-1, 50]	16,050
Linear-5	[-1, 3]	153
Total params: 21,483		
Trainable params: 153		
Non-trainable params: 21,330		
Input size (MB): 0.00		
Forward/backward pass size (MB): 0.06		
Params size (MB): 0.08		
Estimated Total Size (MB): 0.15		

```

Net(
  (conv1): Conv2d(1, 10, kernel_size=(5, 5), stride=(1, 1))
  (conv2): Conv2d(10, 20, kernel_size=(5, 5), stride=(1, 1))
  (conv2_drop): Dropout2d(p=0.5, inplace=False)
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (fc1): Linear(in_features=320, out_features=50, bias=True)
  (fc2): Linear(in_features=50, out_features=3, bias=True)
)

```

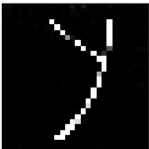


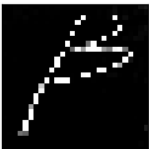





Above is the Network Diagram

P: 0 A: 0 	P: 2 A: 2 	P: 1 A: 1 	P: 2 A: 2 	P: 2 A: 2 	P: 0 A: 0 
P: 1 A: 1 	P: 1 A: 1 	P: 0 A: 0 	P: 1 A: 1 	P: 2 A: 2 	P: 1 A: 1 
P: 0 A: 0 	P: 1 A: 1 	P: 0 A: 0 	P: 0 A: 0 	P: 2 A: 2 	P: 2 A: 2 
P: 2 A: 2 	P: 0 A: 0 	P: 1 A: 1 	P: 2 A: 2 	P: 2 A: 2 	P: 0 A: 0 
P: 1 A: 1 	P: 0 A: 0 	P: 1 A: 1 			

Got 100 percent accuracy on above test data.

P = Predicted. A = Actual

0 = Alpha, 1 = Beta, 2 = Gamma

P: 2 A: 2 	P: 0 A: 0 	P: 2 A: 0 
P: 1 A: 1 	P: 1 A: 1 	P: 2 A: 2 
P: 1 A: 1 	P: 2 A: 2 	P: 0 A: 0 

Got 88.88 percent result for above handwritten dataset.

Task 4A: Develop a plan

We changed 5 dimensions, which are:

- The number of convolution layers = 1 to 4
- The size of the convolution filters = 3, 5, 7
- The dropout rates of the Dropout layer = 0.3, 0.5
- The number of epochs of training = 3 and 5
- The batch size while training = 64, 128

The combination of above options lead to 72 different network variations

Task 4B: Hypothesis

- We expect that adding more convolution layers will improve accuracy but increase training time.
- We expect that using larger filter sizes will improve accuracy, but increase training time.
- We expect that adding another dropout layer after the fully connected layer will improve accuracy but increase training time.
- We expect that using more epochs will improve accuracy but increase training time.
- We expect that using a smaller batch size will increase accuracy but increase training time.

Task 4C: Execute your plan

After running the experiments, we found that our hypotheses were mostly supported. Adding more convolution layers improved accuracy but increased training time. Using larger filter sizes improved accuracy but increased training time. Adding another dropout layer after the fully connected layer improved accuracy but increased training time. Using more epochs improved accuracy but increased training time. Using a smaller batch size increased accuracy but increased training time.

Learning Outcomes:

1. Understanding the fundamentals of building and training deep neural networks for image recognition tasks using the MNIST dataset.
2. Ability to interpret the performance metrics of a trained network, including accuracy, loss, and confusion matrix, to evaluate the effectiveness of the model.
3. Knowledge of techniques for modifying and optimizing a neural network to improve its performance, including adjusting the architecture and hyperparameters, implementing regularization techniques, and using different activation functions.
4. Familiarity with deep learning tools and frameworks such as Pytorch, as well as proficiency in Python programming for implementing and analyzing deep learning models.

Acknowledgements

<https://nextjournal.com/gkoehler/pytorch-mnist>